

Ang Li

☎ 347-952-6772 | ✉ leonli66xyz@gmail.com | 🏠 leonlixyz.github.io | 📺 LeonLixyz

Education

New York University

Ph.D. in Computer Science

Advised by [Pavel Izmailov](#) and [Micah Goldblum](#)

2025 - Present

Columbia University

M.S. in Computer Science

Advised by [Micah Goldblum](#), [Tianyi Peng](#), and [Hongseok Namkoong](#)

2025

Columbia University

B.S. in Computer Science

Magna Cum Laude

2024

Publications

Understanding Reasoning from Pretraining to Post-Training

Arxiv Preprint

Jingyan Shen*, **Ang Li***, Salman Rahman, Yifan Sun, Micah Goldblum, Matus Telgarsky, and Pavel Izmailov

2026

End-to-End Context Compression at Scale

ICML 2026 Workshop on Structured Probabilistic Inference & Generative Modeling

Ang Li*, Sean McLeish*, Haozhe Chen*, Nimit Kalra, Zaiqian Chen, Artem Gazizov, Venkata Anoop Suhas Kumar Morisetty, Bhavya Kailkhura, Harshitha Menon, Zhuang Liu, Brian R. Bartoldson, Tom Goldstein, Sanae Lotfi, Micah Goldblum, Pavel Izmailov

2026

Multi-Token Residual Prediction

arXiv Preprint

Yufeng Xu, Zishuo Bao, Qian Wang, Zeshen Zhang, Haoqi Zhang, Bowen Peng, **Ang Li**, Rahul Chalamala, Yucheng Lu

2026

Measuring the Limits of Continual Learning for LLMs

2nd Workshop on Compositional Learning: Safety, Interpretability, and Agents (CompLearn, ICML 2026)

Nimit Kalra, Narutatsu Ri, Zerzar Bukhari, **Ang Li**, Sanae Lotfi, Liam H Fowl, Micah Goldblum

2026

Teaching Pretrained Language Models to Think Deeper with Retrofitted Recurrence

Conference on Language Modeling (COLM 2026)

Sean McLeish, **Ang Li**, John Kirchenbauer, Dayal Kalra, Brian R. Bartoldson, Bhavya Kailkhura, Avi Schwarzschild, Jonas Geiping, Tom Goldstein, Micah Goldblum

2025

Zebra-CoT: A Dataset for Interleaved Vision Language Reasoning

International Conference on Learning Representations (ICLR 2026)

Ang Li*, Charles Wang*, Deqing Fu*, Kaiyu Yue*, Zikui Cai*, Wang Bill Zhu*, Ollie Liu*, Peng Guo*, Willie Neiswanger, Furong Huang, Tom Goldstein, Micah Goldblum

2025

Twin-2K-500: A dataset for building digital twins of over 2,000 people based on their answers to over 500 questions

Marketing Science

Olivier Toubia, George Z Gui, Tianyi Peng, Daniel J Merlau, **Ang Li**, Haozhe Chen

2025

Commercial LLM Agents Are Already Vulnerable to Simple Yet Dangerous Attacks

arXiv Preprint

Ang Li*, Yin Zhou*, Vethavikashini Chithrra Raghuram*, Tom Goldstein, Micah Goldblum

2025

LLM Generated Persona is a Promise with a Catch

2025

Advances in Neural Information Processing Systems (NeurIPS 2025)

Ang Li, Haozhe Chen, Hongseok Namkoong, Tianyi Peng

Architectural and Inferential Inductive Biases For Exchangeable Sequence Modeling

2025

Advances in Neural Information Processing Systems (NeurIPS 2025)

Daksh Mittal*, Ang Li*, Tzu-Ching Yen*, Daniel Guetta, Hongseok Namkoong

PersonalLLM: Tailoring LLMs to Individual Preferences

2025

International Conference on Learning Representations (ICLR 2025)

Thomas Zollo*, Andrew Siah*, Naimeng Ye, Ang Li, Hongseok Namkoong

QGym: Scalable Simulation and Benchmarking of Queuing Network Controllers

2024

Advances in Neural Information Processing Systems (NeurIPS 2024)

Haozhe Chen*, Ang Li*, Ethan Che*, Tianyi Peng, Jing Dong, Hongseok Namkoong

Experience

Modal Labs

New York, NY

Researcher

Feb 2026 - Present

- Async RL
- Test Time Training and Continual Learning

New York University

New York, NY

Research Assistant

Aug 2025 - Present

- Self Improving Models
- Continual Learning and Long Context
- Science of Pretraining and Post-Training

Relace AI

San Francisco, CA

Research Intern

June 2025 - Aug 2025

- RL for SWE agents

Columbia Engineering, Electrical Engineering

New York, NY

Research Assistant, Advised by Micah Goldblum

Sep 2024 - July 2025

- Visual Chain of Thought
- Agent Safety

Columbia Business School, Decision, Risk, and Operations

New York, NY

Research Assistant, Advised by Tianyi Peng, Jing Dong, and Hongseok Namkoong

Jan 2023 - May 2025

- Efficient RLHF
- LLM Based Simulation of Human
- RL for Queuing system

Columbia's Center for Theoretical Neuroscience at Zuckerman Institute

New York, NY

Research assistant at Professor Larry Abbott's Lab

May 2022 - May 2023

- Memory and Plasticity.

Talks and Presentations

AWS Responsible AI Group

Online

Invited talk

2025

Commercial LLM Agents Are Already Vulnerable to Simple Yet Dangerous Attacks

Columbia's IEOR/DRO Reading Group

NY, New York

Presenter

2023

- **Foundation Models Reading Group:** Delivered a presentation on the scaling laws of large language models
- **Deep Generative Models Reading Group:** Delivered a presentation on the auto-regressive models for visual generation

Teaching

NYU Center for Data Science

Teaching Assistant

- CSCI-GA 3033-125 Special Topics: RL with Foundation Models

NY, New York

Fall 2025 - Present

Columbia Computer Science Department

Teaching Assistant

- COMS 4705 Natural Language Processing (Spring 2023, Fall 2023)

NY, New York

Jan 2023 - Dec 2023

Service

Reviewer

Nature	2025
ICLR	2026
TMLR	2025
NeurIPS	2025
COLM	2025

Skills

Machine Learning Skills	Transformers, Fine-tuning, RLHF, DPO, Deep RL, RL with LLM
Deep Learning Engineering	Pytorch, FSDP, DeepSpeed, Verl
General Engineering	Python, C