

Ang Li

leonli66xyz@gmail.com | leonlixyz.github.io | [GitHub](https://github.com/leonli66xyz) | 347-952-6772

Education

New York University , Ph.D. in Computer Science	2025 – Present
Advised by Pavel Izmailov and Micah Goldblum	
Columbia University , M.S. in Computer Science	2025
Advised by Micah Goldblum , Tianyi Peng , and Hongseok Namkoong	
Columbia University , B.S. in Computer Science	2024
magna cum laude	

Experience

Modal Labs	Feb 2026 – Present
Researcher New York, NY	
- Studied the empirical science of scaling RL, including asynchronous training stability and staleness, long-horizon agent tasks with repeated context compaction, and rollout inference optimization - Developed infrastructure and a benchmark for LLMs to automatically train agents on real-world tasks in law, finance, business, and cybersecurity	
New York University	Aug 2025 – Present
Research Assistant New York, NY	
- Developed an empirical pretraining–RL scaling law across 50+ training configurations spanning model sizes and compute budgets, providing guidance for compute allocation between pretraining and RL posttraining - Scaled encoder-decoder architectures for long-context modeling, training on up to 400B tokens per model and using 600,000+ GPU-hours across all experiments	
Relace AI	Jun 2025 – Aug 2025
Research Intern San Francisco, CA	
Trained LLMs on software engineering tasks with RL, built a training pipeline and sandboxed environments for rollouts and evaluation, and audited agents for reward hacking	
Columbia Engineering	Sep 2024 – Jul 2025
Research Assistant New York, NY	
- Investigated security vulnerabilities in computer-use and scientific discovery agents, including prompt injection, retrieval database poisoning, and safeguard bypasses - Developed a visual chain-of-thought dataset and fine-tuned unified multimodal models for interleaved text-image reasoning	
Columbia Business School	Jan 2023 – May 2025
Research Assistant New York, NY	
- Studied efficient exploration methods for RLHF at scale with uncertainty quantification - Conducted LLM-based simulations across 500+ questions on elections, public opinion, and consumer preferences using over 1M personas and six models, identifying systematic biases introduced by persona generation	
Columbia University, Zuckerman Institute	May 2022 – May 2023
Research Assistant, Center for Theoretical Neuroscience New York, NY	
Studied continual familiarity detection through meta-learned synaptic plasticity and memory addressing, and conducted interpretability analysis of the memory matrix	

Selected Publications

* Equal contribution

Understanding Reasoning from Pretraining to Post-Training	2026
<i>Advances in Neural Information Processing Systems (NeurIPS 2026)</i>	
Jingyan Shen*, Ang Li* , Salman Rahman, Yifan Sun, Micah Goldblum, Matus Telgarsky, and Pavel Izmailov	

End-to-End Context Compression at Scale <i>ICML 2026 Workshop on Structured Probabilistic Inference & Generative Modeling</i> Ang Li* , Sean McLeish*, Haozhe Chen*, Nimit Kalra, Zaiqian Chen, Artem Gazizov, Venkata Anoop Suhas Kumar Morisetty, Bhavya Kailkhura, Harshitha Menon, Zhuang Liu, Brian R. Bartoldson, Tom Goldstein, Sanae Lotfi, Micah Goldblum, and Pavel Izmailov	2026
Teaching Pretrained Language Models to Think Deeper with Retrofitted Recurrence <i>Conference on Language Modeling (COLM 2026)</i> Sean McLeish, Ang Li , John Kirchenbauer, Dayal Singh Kalra, Brian R. Bartoldson, Bhavya Kailkhura, Avi Schwarzschild, Jonas Geiping, Tom Goldstein, and Micah Goldblum	2025
Zebra-CoT: A Dataset for Interleaved Vision Language Reasoning <i>International Conference on Learning Representations (ICLR 2026)</i> Ang Li* , Charles Wang*, Deqing Fu*, Kaiyu Yue*, Zikui Cai*, Wang Bill Zhu*, Ollie Liu*, Peng Guo*, Willie Neiswanger, Furong Huang, Tom Goldstein, and Micah Goldblum	2025
Twin-2K-500: A dataset for building digital twins of over 2,000 people based on their answers to over 500 questions <i>Marketing Science</i> Olivier Toubia, George Z. Gui, Tianyi Peng, Daniel J. Merlau, Ang Li , and Haozhe Chen	2025
Commercial LLM Agents Are Already Vulnerable to Simple Yet Dangerous Attacks <i>arXiv Preprint</i> Ang Li* , Yin Zhou*, Vethavikashini Chithrara Raghuram*, Tom Goldstein, and Micah Goldblum	2025
LLM Generated Persona is a Promise with a Catch <i>Advances in Neural Information Processing Systems (NeurIPS 2025)</i> Ang Li , Haozhe Chen, Hongseok Namkoong, and Tianyi Peng	2025
PersonalLLM: Tailoring LLMs to Individual Preferences <i>International Conference on Learning Representations (ICLR 2025)</i> Thomas Zollo*, Andrew Siah*, Naimeng Ye, Ang Li , and Hongseok Namkoong	2025

Talks & Service

AWS Responsible AI Group , Invited Talk AI Agent Safety	2025 Online
Columbia DRO Group , Invited Talk Reinforcement Learning for LLM Reasoning	2025 New York, NY
NYU Center for Data Science , Teaching Assistant CSCI-GA 3033-125 Special Topics: RL with Foundation Models	Fall 2025 – Dec 2025 New York, NY
Columbia Computer Science Department , Teaching Assistant COMS 4705 Natural Language Processing (Spring 2023, Fall 2023)	Jan 2023 – Dec 2023 New York, NY
Reviewer: Nature (2025), NeurIPS (2025, 2026), ICLR (2026, 2027), COLM (2025), TMLR (2025)	

Skills

ML Engineering: PyTorch, FSDP, DeepSpeed, Megatron, TorchTitan, Miles, slime, verl